

Data Science Capstone Project

April 23, 2021

# Cutting Server Costs with Machine Learning

Maritsa Norton & Gianluca Macauda



{**Pro**pulsion}

# The team



**MARITSA NORTON**

- 7 years' strategic analytics in tech (Expedia / OpenTable)
- Background in economics



**GIANLUCA MACAUDA**

- Spent last 7 years in research
- Background in cognitive neuroscience and psychology

# Introduction to CodeNotary


Many companies rely on advanced server platforms to run their business applications.

CodeNotary provides immutable database management and monitoring services that sit on top of these server platforms.

Fig. 1: Simplified NUMA server architecture

## Servers

the physical machines that manage the VMs



## Virtual machines (VMs)

Run business applications for this type of server architecture

Default settings can lead to **misconfigurations** of the server settings and thus **inefficiency**.

# Opportunity to help clients reduce server costs



## Undiagnosed server **misconfiguration**


- Slow applications
- High costs (adds approx. 20% to budget )

# Opportunity to help clients reduce server costs



Undiagnosed server  
**misconfiguration**

- Slow applications
- High costs (adds approx. 20% to budget )



Option 1: **Buy more servers**

- Expensive
- May not solve root problem

# Opportunity to help clients reduce server costs



## Undiagnosed server **misconfiguration**

- Slow applications
- High costs (adds approx. 20% to budget )

### Option 1: Buy more servers

- Expensive
- May not solve root problem

### Option 2: **Improve server efficiency**

- + Lower costs in the long term
- Requires downtime and expertise
- Cannot quantify expected outcomes

# Opportunity to help clients reduce server costs



## Undiagnosed server **misconfiguration**

- Slow applications
- High costs (adds approx. 20% to budget )

### Option 1: Buy more servers

- Expensive
- May not solve root problem

### Option 2: **Improve server efficiency**


- + Lower costs in the long term
- Requires downtime and expertise
- ~~- Cannot quantify expected outcomes~~

# The goal & the approach:

Machine Learning to forecast server inefficiency




## Datasource



**Configuration** settings and **performance** measures for:

5,406 **Virtual machines (VM)** on  
230 **Servers**

## Efficiency Forecasting: Predictive Modeling




### Forecast individual VM efficiency

Capture the dynamics of the system

### Forecast aggregate server efficiency

Quantify the business impact of specific configuration changes

## Application




1. **Quantify** the impact of **configuration changes**



# Modeling & evaluation:

Predicting each individual virtual machine's weekly efficiency



## Model insights


- + Model can **capture** most of the patterns
- **Underestimation** of real inefficiency

## Proof of concept

Can we use these modeling techniques  
to predict efficiency gains after a  
configuration change?

# Proof of concept:

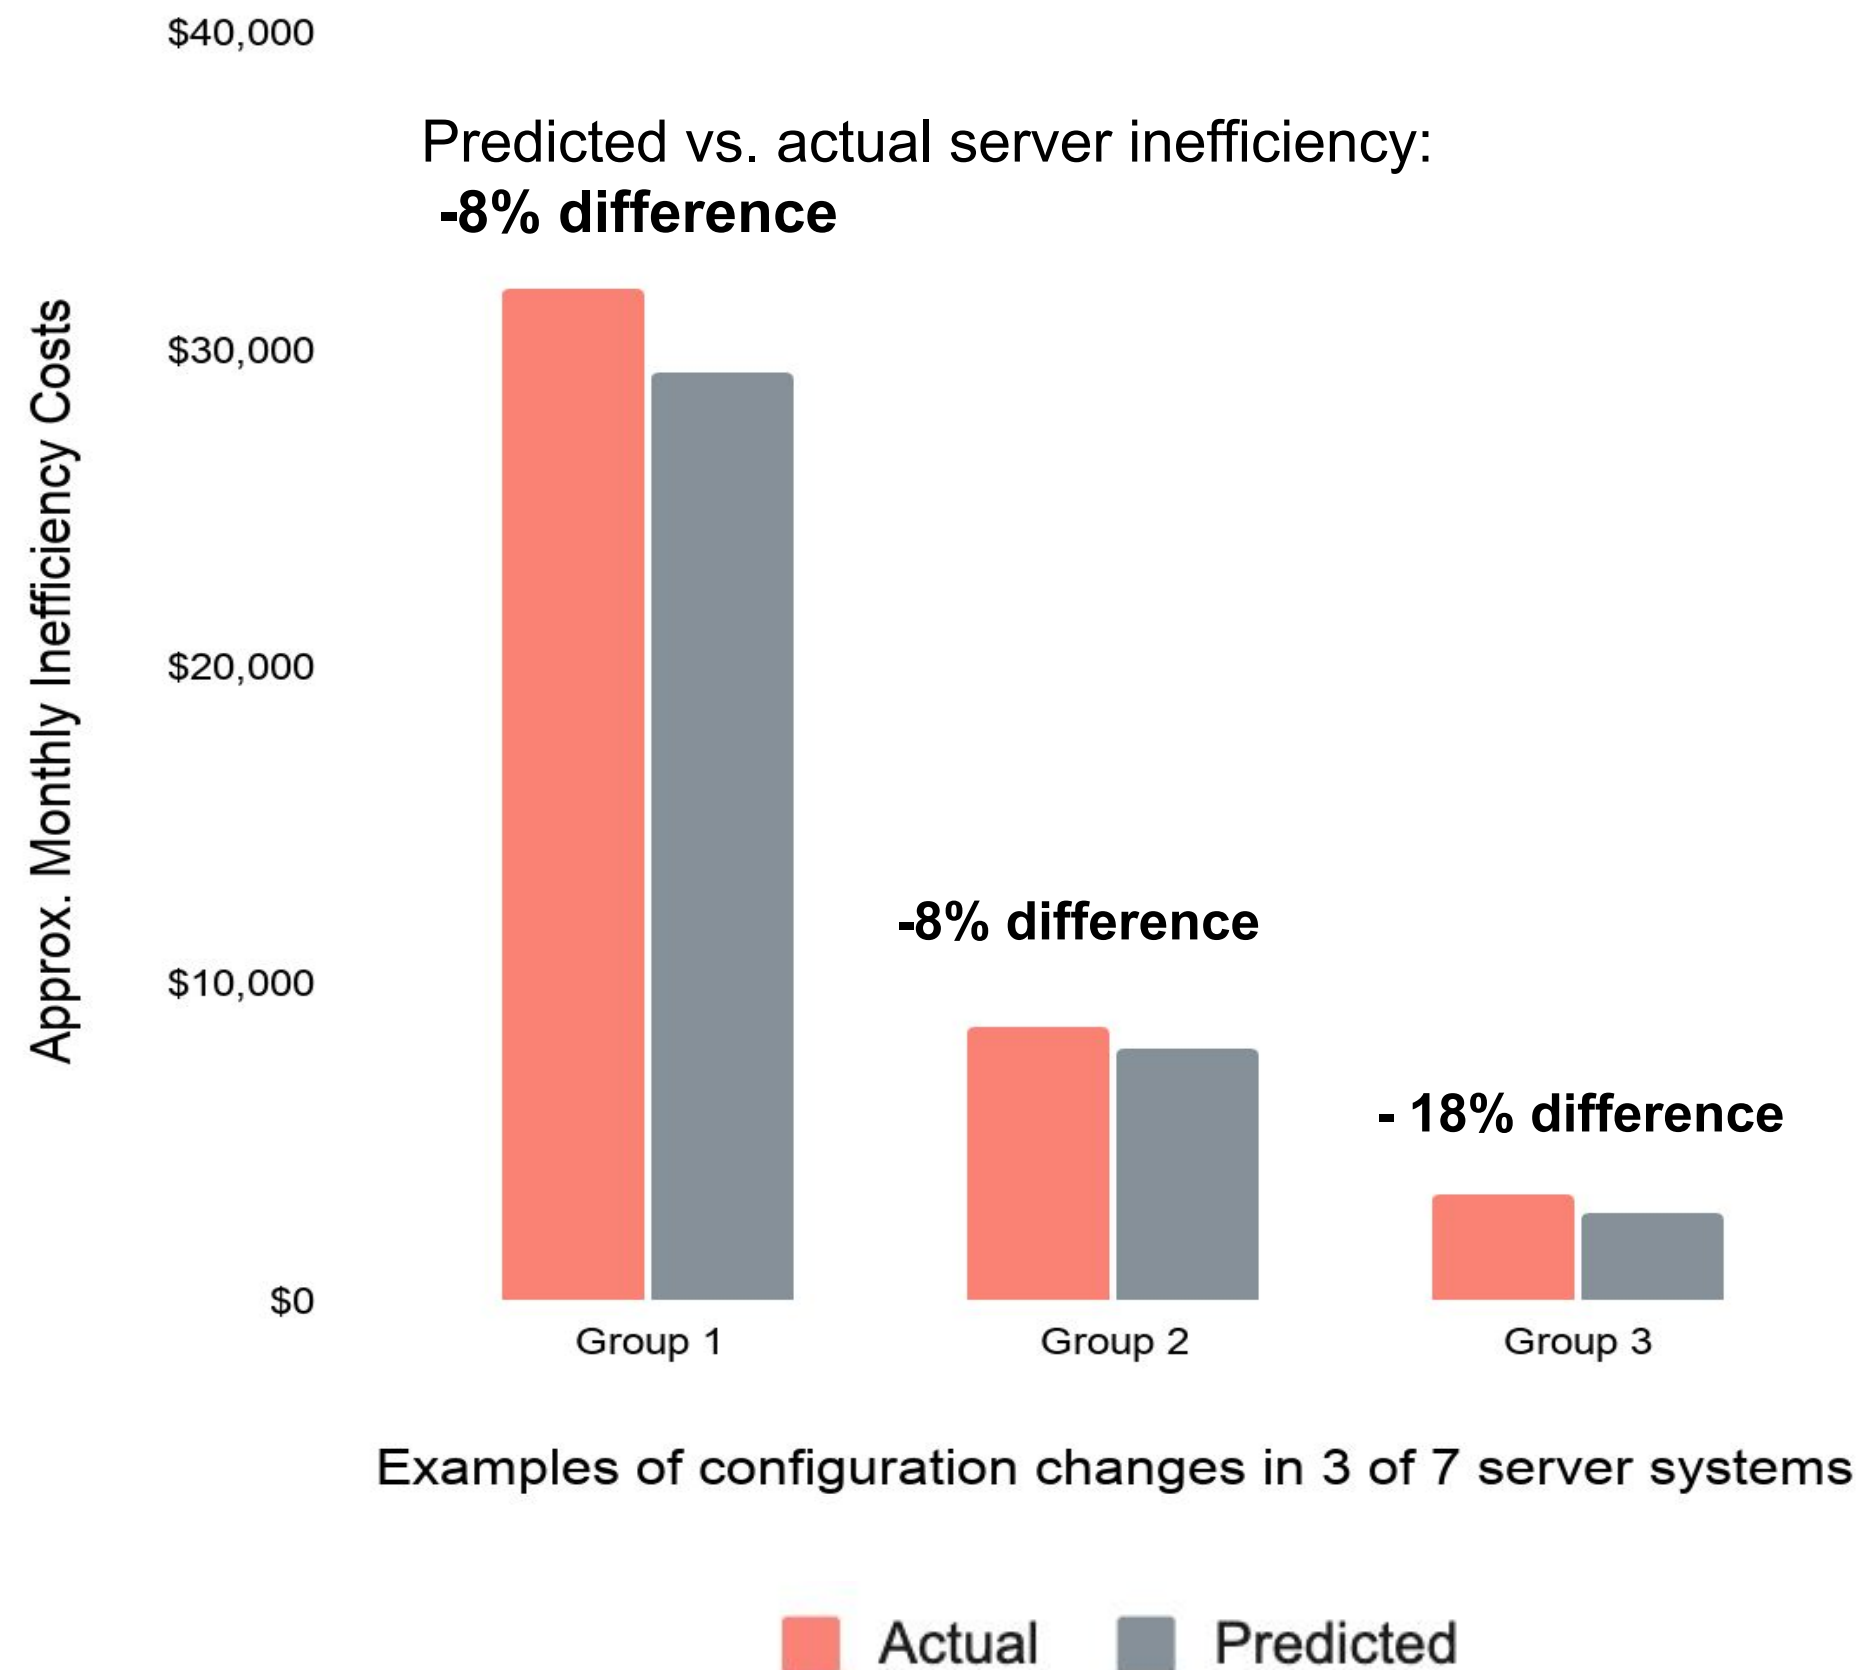
ML model to predict \$\$ impact of simple configuration changes



- + Predictions underestimate the system's actual inefficiency by -8 to -18%
- + **Supports a cost-benefit assessment** of proposed configuration changes (for VM-level changes to poorly configured platforms)

# Proof of concept:

ML model to predict \$\$ impact of simple configuration changes



- + Predictions underestimate the system's actual inefficiency by -8 to -18%
- + **Supports a cost-benefit assessment** of proposed configuration changes (for VM-level changes to poorly configured platforms)
- Room to improve accuracy
- Only useful for VM configuration changes on very inefficient systems

# Takeaways



## Link VMs to servers

- Connection of VM and server settings are critical
- Current stage: No possibility to connect them in satisfactory way

## Imbalanced dataset leads to a “biased” model


- One client had no inefficiency problems (Average Inefficiency = 1.19%)
- Clients with efficiency problems & more sophisticated modeling approaches

## Promising efficiency forecasting (3 clients)

- mean absolute error around 1%
- For **inefficient VMs** (VMs of interest): average **underestimation of 6%**


# Next Steps & Outlook

## Data collection




1. Diverse data
2. Link Servers to VMs

## Modeling



Capture dynamics of servers and VMs

## Application




Case-specific predictions  
Dynamic simulations of  
configuration changes




# Next Steps & Outlook

## Data collection




1. Diverse data
2. Link Servers to VMs

## Modeling



Capture dynamics of servers and VMs

## Application



Case-specific predictions  
Dynamic simulations of  
configuration changes

Novel tool to help organizations stop wasting money

# Thank you



**MARITSA NORTON**



[maritsa.norton.oleson@gmail.com](mailto:maritsa.norton.oleson@gmail.com)



[www.linkedin.com/in/maritsa-norton/](https://www.linkedin.com/in/maritsa-norton/)



**GIANLUCA MACAUDA**



[gianluca@macauda.ch](mailto:gianluca@macauda.ch)



[www.linkedin.com/in/gianluca-macauda/](https://www.linkedin.com/in/gianluca-macauda/)